

▼ Lab Activity - 1

1. Implement a Linear Regression model in Machine Learning and fit the model to predict total vaccination by 15-Mar-2021 for India country.

Using the country_vaccinations dataset.

```
import pandas as pd
from datetime import datetime
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
sns.set_style("whitegrid")
```

```
covid_df = pd.read_csv("/content/drive/MyDrive/Colab Notebooks/Datasets/covid_vaccinations.csv")
print(covid_df.shape)
covid_df.head()
```

```
(4435, 15)
```

	country	iso_code	date	total_vaccinations	people_vaccinated	people_fullly_vaccinated
0	Albania	ALB	2021-01-10	0.0	0.0	0.0
1	Albania	ALB	2021-01-11	NaN	NaN	NaN
2	Albania	ALB	2021-01-12	128.0	128.0	128.0
3	Albania	ALB	2021-01-13	188.0	188.0	188.0
4	Albania	ALB	2021-01-14	266.0	266.0	266.0

```
covid_df.isnull().sum()
```

```
country      0
iso_code    304
```

```

date                                0
total_vaccinations                  1519
people_vaccinated                    1952
people_fully_vaccinated              2773
daily_vaccinations_raw               1968
daily_vaccinations                   154
total_vaccinations_per_hundred       1519
people_vaccinated_per_hundred        1952
people_fully_vaccinated_per_hundred  2773
daily_vaccinations_per_million       154
vaccines                             0
source_name                          0
source_website                       0
dtype: int64

```

```
covid_df['date'].min(), covid_df['date'].max()
```

```
('2020-12-08', '2021-02-27')
```

```
covid_df['country'].value_counts()
```

```

Lithuania      82
Scotland       76
United Kingdom  76
England         76
Wales           76
..
Senegal         5
South Korea     3
Ukraine         3
Saint Helena    1
Greenland       1
Name: country, Length: 112, dtype: int64

```

```

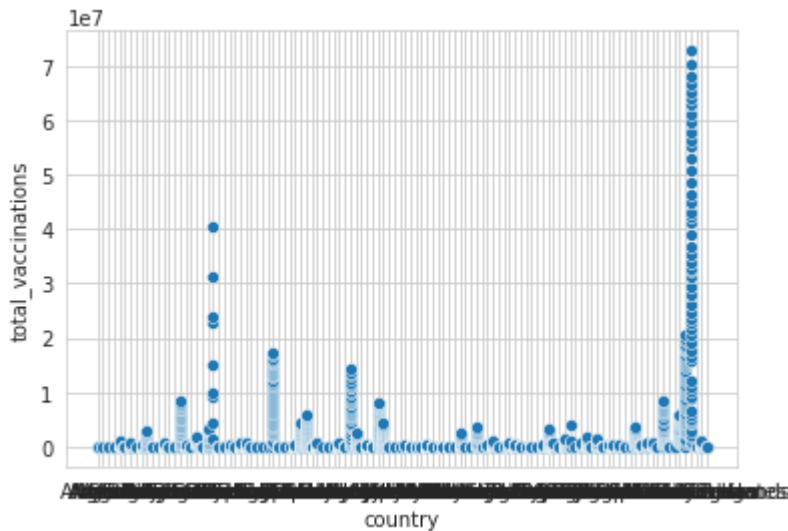
covid_req_data = covid_df.loc[:, ['country', 'date', 'total_vaccinations']]
covid_req_data['date'] = pd.to_datetime(covid_req_data['date'])
covid_req_data.head(10)

```

	country	date	total_vaccinations
0	Albania	2021-01-10	0.0
1	Albania	2021-01-11	NaN


```
sns.scatterplot(x='country', y='total_vaccinations', data=covid_req_data)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f5ae9ed5210>
```



```
sns.scatterplot(x='date', y='total_vaccinations', data=covid_req_data)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f5ae9d03a90>
```



```
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
```

```
x1 = le.fit_transform(covid_req_data['country'])
x2 = le.fit_transform(covid_req_data['date'])
X = np.array([x1, x2]).T
covid_req_data['total_vaccinations'].fillna(0, inplace=True)
```

```
covid_req_data['total_vaccinations'].fillna(0, inplace=True)
y = covid_req_data['total_vaccinations'].values
```

```
x = covid_req_data.groupby(['date']).sum()
x['date_diff']=x['total_vaccinations']
count=0
for index, row in x.iterrows():
    row['date_diff']=count
    count+=1
```

```
x = x[['date_diff']]
y = x['total_vaccinations']
```

```
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, r2_score, accuracy_score
```

```
X_train, X_test, y_train, y_test = train_test_split(x, y, test_size=0.33, r
```

```
rg = LinearRegression()
rg.fit(X_train, y_train)
y_pred = rg.predict(X_test)
```

```
print(f'Correlation Score: {rg.score(X_test, y_test)} \nMSE: {mean_squared
```

```
Correlation Score: 0.8332154888981802
MSE: 864569548661889.8
```

2. Implement a Multiple Regression model in Machine Learning to fit

- ▼ the model. You can assume features / independent variables and dependent variable.

Use winequality-red.csv dataset.


```
wine = pd.read_csv("/content/drive/MyDrive/Colab Notebooks/Datasets/winequ
print(wine.shape)
wine.head()
```

(1599, 12)

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	
0	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.
1	7.8	0.88	0.00	2.6	0.098	25.0	67.0	0.9968	3.

sns.pairplot(wine)

<seaborn.axisgrid.PairGrid at 0x7f5ae8530f90>



```
x = wine.iloc[:, :-1].values
y = wine.loc[:, ['quality']].values
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25)
```

► Multiple Linear Regression

```
from sklearn.linear_model import LinearRegression
```

```

rg = LinearRegression()

rg.fit(X_train, y_train)
y_pred = np.round(rg.predict(X_test))

print(f'Correlation Score: {r2_score(y_test, y_pred)} \nAccuracy Score: {a

Correlation Score: 0.29698656219105857
Accuracy Score: 0.6025
MSE: 0.48

```

▼ Polynomial Regression

```

from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import PolynomialFeatures
poly_reg = PolynomialFeatures(degree = 2)
x_poly = poly_reg.fit_transform(X_train)
linear_reg = LinearRegression()
linear_reg.fit(x_poly, y_train)
y_pred = np.round(linear_reg.predict(poly_reg.transform(X_test)))

print(f'Correlation Score: {r2_score(y_test, y_pred)} \nAccuracy Score: {a

Correlation Score: 0.10658708945113693
Accuracy Score: 0.6075
MSE: 0.61

```

▼ 4. Implement a Decision Tree model in Machine Learning and fit the model towards the regression to predict housing price.

Use melb_data.csv dataset.

```

house = pd.read_csv("/content/drive/MyDrive/Colab Notebooks/Datasets/melb_
print(house.shape)
house.head()

```

(13580, 21)

	Suburb	Address	Rooms	Type	Price	Method	SellerG	Date	Distance
0	Abbotsford	85 Turner St	2	h	1480000.0	S	Biggin	3/12/2016	2.5
1	Abbotsford	25 Bloomburg St	2	h	1035000.0	S	Biggin	4/02/2016	2.5
2	Abbotsford	5 Charles St	3	h	1465000.0	SP	Biggin	4/03/2017	2.5

```
house.isnull().sum()
```

```
Suburb          0
Address         0
Rooms          0
Type            0
Price           0
Method          0
SellerG         0
Date            0
Distance        0
Postcode        0
Bedroom2        0
Bathroom        0
Car             62
Landsize        0
BuildingArea    6450
YearBuilt       5375
CouncilArea     1369
Latitude        0
Longitude       0
Regionname      0
Propertycount   0
dtype: int64
```

```
values = {'BuildingArea': house['BuildingArea'].mean(), 'YearBuilt': house
house.fillna(value=values, inplace=True)
```

```
house.dropna(inplace=True)
house.drop(['Date'], axis=1, inplace=True)
```

```
objList = house.select_dtypes(include = "object").columns
print (objList)
```

```
Index(['Suburb', 'Address', 'Type', 'Method', 'SellerG', 'CouncilArea',
      'Regionname'],
      dtype='object')
```

```
#Label Encoding for object to numeric conversion
from sklearn.preprocessing import LabelEncoder
```



```
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
```

```
for feat in objList:
    house[feat] = le.fit_transform(house[feat].astype(str))
```

```
house.head(10)
```

	Suburb	Address	Rooms	Type	Price	Method	SellerG	Distance	Postcode
0	0	2817	2	0	1480000.0	1	41	2.5	3067.0
1	0	8590	2	0	1035000.0	1	41	2.5	3067.0
2	0	843	3	0	1465000.0	3	41	2.5	3067.0
3	0	11974	3	0	850000.0	0	41	2.5	3067.0
4	0	1717	4	0	1600000.0	4	202	2.5	3067.0
5	0	4442	2	0	941000.0	1	198	2.5	3067.0
6	0	4384	3	0	1876000.0	1	202	2.5	3067.0
7	0	3398	2	0	1636000.0	1	202	2.5	3067.0
8	0	2251	1	2	300000.0	1	41	2.5	3067.0
9	0	3183	2	0	1097000.0	1	41	2.5	3067.0

```
x = house.drop(['Price'], axis=1)
y = house.loc[:, ['Price']]
X_train, X_test, y_train, y_test = train_test_split(x, y, test_size=0.25)
```

```
from sklearn.tree import DecisionTreeRegressor
dr = DecisionTreeRegressor()
```

```
dr.fit(X_train, y_train)
y_pred = np.round(dr.predict(X_test))
```

```
print(f'Correlation score: {r2_score(y_test, y_pred)} \nMSE: {mean_squared_error(y_test, y_pred)}')
```

```
Correlation score: 0.5482045701918146
MSE: 187114525545.24075
```

5. Implement a Logistic Regression model in Machine Learning and fit the model to predict heart rate of a person based on age and BMI.

Use Framingham.csv dataset.

```
data = pd.read_csv("/content/drive/MyDrive/Colab Notebooks/Datasets/framingham.csv")
print(data.shape)
data.head()
```

```
(4240, 16)
```

	male	age	education	currentSmoker	cigsPerDay	BPMeds	prevalentStroke	prevalentHyp	diabetes	totChol	sysBP	diaBP	BMI	heartRate	glucose	TenYearCHD
0	1	39	4.0	0	0.0	0.0	0	0	0	162.0	101	54	25.25	67	99	0
1	0	46	2.0	0	0.0	0.0	0	0	0	162.0	101	54	25.25	67	99	0
2	1	48	1.0	1	20.0	0.0	0	0	0	162.0	101	54	25.25	67	99	0
3	0	61	3.0	1	30.0	0.0	0	0	0	162.0	101	54	25.25	67	99	0
4	0	46	3.0	1	23.0	0.0	0	0	0	162.0	101	54	25.25	67	99	0

```
data["glucose"].fillna((data["glucose"].mode())[0], inplace=True)
data.dropna(inplace=True)
data.isnull().sum()
```

```
male          0
age           0
education     0
currentSmoker 0
cigsPerDay    0
BPMeds        0
prevalentStroke 0
prevalentHyp  0
diabetes      0
totChol       0
sysBP         0
diaBP         0
BMI           0
heartRate     0
glucose       0
TenYearCHD    0
dtype: int64
```

```
target1=data[data['TenYearCHD']==1]
target0=data[data['TenYearCHD']==0]
```

```
target1=resample(target1,replace=True,n_samples=len(target0),random_state=42)
```

```
target=pd.concat([target0,target1])
```

```
target['TenYearCHD'].value_counts()
```

```
data=target
```

```
data = img  
np.shape(data)
```